

Backpropagation

Hikmat Farhat

November 18, 2020

1 Supervised Learning of Classification Problems

We derive the formulas needed to perform backpropagation in a fully connected neural network when the problem is supervised learning of k classes. Since we have multiple *distinct* classes the labels are represented in one_hot encoding. So $y = [10 \dots 0]$ represents class 1, $y = [01 \dots 0]$ represents class 2, etc.

Let \mathbf{x} be the input, our goal is to compute $p(y | \mathbf{x})$. Following Bishop we write

$$p(y | \mathbf{x}) = \prod_{i=1}^k p_i^{y_i} \quad (1)$$

Where p_i is the probability that $y = c_i$ given that the input is x , y is the one_hot representation of the class. This means that only one of the y_i is 1 and the rest are zeros with $\sum_i p_i = 1$. This implies

$$p(y = y_i | \mathbf{x}) = p_i$$

A special case is when we have only two categories then $p_1 = 1 - p_0$ and we can write

$$p(y | \mathbf{x}) = p_0^y \cdot (1 - p_0)^{1-y}$$

2 Kullback-Leibler

In classification problem we have a neural network that approximates $p(y | \mathbf{x})$. Let z be the output of the neural network then we model the probabilities by

$$\hat{p}_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

The loss is defined by the Kullback-Leibler divergence

$$\begin{aligned} KL(p \parallel \hat{p}) &= \sum p \log \frac{\hat{p}}{p} \\ &= \sum p \log \hat{p} - p \log p \geq 0 \end{aligned} \quad (2)$$

Since we can't influence the unknown distribution p we minimize the first term, or alternatively we maximize its negative (likelihood). The summation in equation (2) is an average over the true (unknown) distribution p . We approximate by averaging over the samples. Replacing the expression in (1) we define the *loss* function as:

$$\mathcal{L} = - \sum_s \sum_k y_{sk} \log \hat{p}_{sk}$$

Where s is the index of the sample and y_{sk} is the one_hot vector encoding of the label of sample s . Note that for each s there is only one $y_{sk} = 1$ and the rest are zeros.

3 Gradient

The loss function \mathcal{L} depends on all the parameters of the neural network. We will see that computing the derivatives with respect to the parameters of one layer depend on the derivatives of the **next** layer.

The output of layer l , is a vector a^l given by $a_{sj}^l = f(z_{sj}^l)$ where f is some nonlinear function, called the activation function, s is the index of the sample, j is the component of a and z , and

$$z_{sj}^{l+1} = \sum_k a_{sk}^l w_{kj}^l + b_j^l \quad (3)$$

In a neural network with $L + 1$ layers, used for classification problems, we usually use the *softmax* function as the activation of the last layer a^L . Furthermore, we interpret the output of the last layer as the probability of a given class, i.e. $a^L \equiv \hat{p}$

3.1 Derivation with respect to the last parameters

Let L be the index of the last layer so $a^L = \sigma(z^L)$ where σ is the softmax function (or sigmoid for 2 classes).

Since the samples are independent we will drop the sample index in what follows for simplicity. To compute the gradient we need an expression for the following quantity:

$$\frac{\partial \mathcal{L}}{\partial z_j^l}$$

3.2 Derivative of the last layer

To simplify the notation we use $z \equiv z^L$ and $a \equiv a^L$ and we drop the index of the samples (i.e. we consider one sample. All the others are the same)

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial z_j} &= -\frac{\partial}{\partial z_j} \sum_k y_k \log a_k \\ &= -\sum_k \frac{y_k}{a_k} \frac{\partial a_k}{\partial z_j}\end{aligned}$$

Since

$$a_k = \frac{e^{z_k}}{\sum_t e^{z_t}}$$

then

$$\frac{\partial a_k}{\partial z_j} = \frac{\delta_{ik} \sum_t e^{z_t} - e^{z_k} e^{z_j}}{(\sum_t e^{z_t})^2}$$

Where δ_{ik} is the Kronecker delta. Multiplying by $\frac{1}{a_k}$ we get

$$\frac{1}{a_k} \frac{\partial a_k}{\partial z_j} = \delta_{jk} - a_j$$

Therefore

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial z_j} &= \sum_k (a_j y_k - \delta_{jk} y_k) \\ &= a_j - y_j\end{aligned}$$

Using Δ_j^L for $\frac{\partial \mathcal{L}}{\partial z_j}$ and reinserting the sample index we get

$$\Delta_{sj}^L = a_{sj} - y_{sj}$$

3.3 Backpropagation

In this section we derive an expression for Δ^l in terms of Δ^{l+1} .

$$\begin{aligned}\Delta_j^l &= \frac{\partial \mathcal{L}}{\partial z_j^l} = \sum_k \frac{\partial \mathcal{L}}{\partial z_k^{l+1}} \frac{\partial z_k^{l+1}}{\partial z_j^l} \\ &= \sum_k \Delta_k^{l+1} \frac{\partial z_k^{l+1}}{\partial a_j^l} \frac{\partial a_j^l}{\partial z_j^l} \\ &= \sum_k \Delta_k^{l+1} w_{jk}^l \theta_j^l\end{aligned}$$

Where $\theta_j^l \equiv \frac{\partial a_j^l}{\partial z_j^l}$ is the derivative of the activation function of layer l with respect to its parameter. Reinserting the sample index we get

$$\Delta_{sj}^l = \sum_k \Delta_{sk}^{l+1} w_{jk}^l \theta_{sj}^l$$

3.3.1 Gradient wrt parameters

Since the neural network is "optimized" by varying the parameters w^l, b^l we need an expression for the gradient of the loss wrt to the parameters.

$$\frac{\partial \mathcal{L}}{\partial w_{ij}^l} = \frac{\partial \mathcal{L}}{\partial z_j^{l+1}} \frac{\partial z_j^{l+1}}{\partial w_{ij}^l}$$

Using eq. (??) we get

$$\frac{\partial \mathcal{L}}{\partial w_{ij}^l} = \Delta_j^{l+1} a_i^l$$

Reinserting the sample index and averaging over the samples (total m samples):

$$\frac{\partial \mathcal{L}}{\partial w_{ij}^l} = \frac{1}{m} \sum_s a_{si}^l \Delta_{sj}^{l+1}$$

Similarly for the bias we get

$$\begin{aligned} \frac{\partial \mathcal{L}}{\partial b_j^l} &= \frac{\partial \mathcal{L}}{\partial z_j^{l+1}} \frac{\partial z_j^{l+1}}{\partial b_j^l} \\ &= \Delta_j^{l+1} \end{aligned}$$

Averaging over the m samples

$$\frac{\partial \mathcal{L}}{\partial b_j^l} = \frac{1}{m} \sum_s \Delta_{sj}^{l+1}$$